

# Bases de données spécialisées

## A CQL(SH) primer



2025-2026

Giovanni Bernardi, [gioXYZirif.fr](http://gioXYZirif.fr)

<http://www.irif.fr/~gio/index.xhtml>

## Warm-up

Why weak consistency ?

## Warm-up

### Why weak consistency ?

#### Theorem (CAP)

*It is impossible in the asynchronous network model to implement a read/write data object that guarantees the following properties:*

- ▶ Availability
- ▶ Atomic Consistency

*in all fair executions (including those in which messages are lost).*

published first at PODC 2002

## Warm-up



is designed to offer which properties out of

Strong Consistency,    Availability,    Partition tolerance

?

## Warm-up



is designed to offer which properties out of

Strong Consistency,    **Availability**,    **Partition tolerance**

## Thus far

- ▶ CAP theorem
- ▶ Anomalies and weak consistency
  - Write Skew example in PostgreSQL
- ▶ Cassandra
  - outline of the distributed system
  - installation, configuration
  - sketch of node management

# Questions we will address

Server side

**How to set-up and run a cassandra node ?**

Last week

Client side

**How to interact with the storage system ?**

## Server side



## Client side



# How to connect to a Cassandra server ?

cqlsh

synopsis simplified:

```
$ cqlsh [addr_ip] [port] [-u username] [-p password]  
      [-e query]
```

basic usage to connect to localhost

```
$ cqlsh
```

options can be given in `~/.cassandra/cqlshrc`

# SHOW

Directive to make cqlsh display information about

**VERSION:** version of cqlsh, Cassandra server,  
CQL specification, and native protocol

Usually displayed as first line at connection time

**HOST:** the node details for host of the cqlsh session

**SESSION tr\_session\_id:** activity details for a query identified  
by `tr_session_id`. Session ids are stored in the table  
`system_traces.sessions`.

## TRACING ON | OFF

Directive to make cqlsh (en|dis)able storage of metadata for all CQL statements in the current session.

Try executing the following queries

```
tracing off
```

```
select * from keyspace_name.table_name;
```

```
tracing on
```

```
select * from keyspace_name.table_name;
```

where `keyspace_name`, `table_name` are names of entities that exist in the cluster.

try running: `select * from system_traces.sessions;`

- ▶ Useful to troubleshoot performance issues
- ▶ Documentation

## PAGING ON | OFF

Makes cqlsh (en|dis)able  
displaying results in chunks of 100 line.

Try executing the following queries

```
paging off
```

```
select * from system.compaction_history;
```

```
paging on
```

```
select * from system.compaction_history;
```

- ▶ Documentation

CLEAR

CTRL + I

## DESCRIBE

Directive to make cqlsh display information about

**CLUSTER**: the cluster name, partitioner, and token ranges.

**SCHEMA**: the schema (i.e. the definitions) of all non-system keyspaces.

**FULL SCHEMA**: also the definitions of the system keyspaces.  
`system, system_traces, system_distributed, ...`

## DESCRIBE

Makes cqlsh display information about

**TABLES:** Every keyspace name, and for each keyspace the names of all the tables it contains.

**TABLE ks\_name.table\_name:** The schema of the table with name `table_name` in the keyspace with name `ks_name`

Try executing the following queries

```
describe tables
```

```
describe system_traces.sessions
```

## USE keyspace\_name

Makes the session “connect” to the given keyspace.

`keyspace_name`: This parameter is the name of the keyspace to connect to. Use quotes ("keyspace\_name") for case-sensitive syntax.

It is the analogous of the instruction `\c dbname` in psql.

## CONSISTENCY

Directive to make cqlsh show or modify the consistency level *within one session!*

no options: Display the current consistency level

`consistency_level` : Set the consistency level to `consistency_level`. As expected, the consistency level determines data availability versus data accuracy for transactions during a session.

## Handy manner to use cqlsh

```
$ cqlsh -e "your_query"
```

Useful to write scripts that query the storage system.

### Example

```
$cqlsh -e "describe keyspaces"
```

Of course the point of writing scripts is to manipulate the output of queries

```
cqlsh -e "select * from system_schema.keyspaces;" | grep SimpleStrategy
```

## Python driver

```
pip3 install cassandra-driver
```

# Cassandra Query Language

looks similar to SQL

but it behaves differently !



# KEYSPACE

Operations: CREATE, **USE** not sql, ALERT, DROP

## Example

```
CREATE KEYSPACE [IF NOT EXISTS] demo
WITH replication = { 'class' : 'NetworkTopologyStrategy',
                     'datacenter1' : 3
                     'datacenter2' : 2 };
```

- ▶ Official documentation

## Replication class


### SimpleStrategy

- ▶ [...] returns the RF nodes that lie right next to each other on the ring.
- ▶ “generally not a wise choice for production because it [...] does not respect datacenter layouts and can lead to wildly varying query latency.”

### NetworkTopologyStrategy

- ▶ “allows to set the replication factor independently for each data-center.”
- ▶ `cassandra-rackdc.properties`

# SimpleStrategy




A client asks Cassandra to write the row (69000, Lyon) in a keyspace with replication factor 3.

The server

1. computes the token,  $\text{hash}(69000) = 67$
2. finds the node responsible for token 67, Node 4
3. writes  $(69000, \text{Lyon})$  in Node 4
4. and in Node 3 and Node 2

# SimpleStrategy




A client asks Cassandra to write the row (69000, Lyon) in a keyspace with replication factor 3.

The server

1. computes the token,  $\text{hash}(69000) = 67$
2. finds the node responsible for token 67, Node 4
3. writes  $(69000, \text{Lyon})$  in Node 4
4. and in Node 3 and Node 2

# Network Topology Strategy




A client asks Cassandra to write the row (69000, Lyon) in a keyspace with replication factor 4 ( $2 + 2$ ).

The server

1. computes the token,  $\text{hash}(69000) = 67$
2. finds the node responsible for token 67, Node 4
3. writes  $(69000, \text{Lyon})$  in Node 4
4. and in enough nodes in different dc/racks (example: 1, 5, 6)

# Network Topology Strategy



A client asks Cassandra to write the row (69000, Lyon) in a keyspace with replication factor 4 ( $2 + 2$ ).

The server

1. computes the token,  $\text{hash}(69000) = 67$
2. finds the node responsible for token 67, Node 4
3. writes  $(69000, \text{Lyon})$  in Node 4
4. and in enough nodes in different dc/racks (example: 1, 5, 6)

# TABLE

Operations: CREATE, ALTER, **TRUNCATE** not sql, DROP

## Example

```
CREATE TABLE [IF NOT EXISTS] mytable (
    cname1 type1,
    :
    cnameN typeN,
    PRIMARY KEY (cnameA, ..., cnameM))
                                         partitionK
    WITH comment='2 buffalo buffalo buffalo buffalo'
```

- ▶ First column in PK identifies **partition**
- ▶ In the primary key it is crucial to have columns to find a row within a partition clustering columns
- ▶ Official documentation

# TABLE

```
CREATE TABLE simplepk (
    id int,
    title text,
    band text,
    PRIMARY KEY (id, title, band))
```

```
CREATE TABLE compositepk (
    id int,
    title text,
    band text,
    PRIMARY KEY ((id, title), band))
```

Different syntax, different meaning:

- ▶ **id** partition key;                    clustering columns: title, band
- ▶ **(id, title)** composite partition key;    clustering column: band

# TABLE

## Example

```
TRUNCATE TABLE dummy;
```

- ▶ TRUNCATE deletes all the data in a table, but not the table
- ▶ Alternative to DROP TABLE

# INSERT

```
insert into  
tablename(cname1,cname2,cname3, ...)  
values (v1,v2,v3, ...);
```

## Example

```
insert into dummy(id,num) values (0,10);  
insert into dummy(id) values (0);
```

- ▶ requires the entire primary key, but not the other values
- ▶ missing values are set to null
- ▶ **it is not SQL:** if the same primary key is used in two inserts  
the row is overwritten!
- ▶ Official documentation

Semantics of SQL and CQL inserts are different!

ex0-insert-refintegrity.sh

ex0-insert-refintegrity.py

# UPDATE

## Example

```
update dummy set num=54 where id=9;
```

```
update dummy set num=1 where id=33;
```

## UPSERT

no fundamental difference between insert and update operations

- ▶ update a row and the primary key does not exist  
⇒ Cassandra creates the row with exceptions
- ▶ insert a row w/ duplicate primary key  
⇒ Cassandra replaces the row
- ▶ Official documentation

`ex1-update.py`

# Questions ?

# Questions ?



J. W. Waterhouse, 1896

## Consider the following table schema

```
create table  
dots(id int, color text, primary key (id, color))  
with  
clustering order by (color desc)
```

- ▶ The partition key is id
- ▶ **Clustering order:** The order of the rows within each partition depends on color

# SELECT

## Example

```
select table_name,comment from system_schema.tables;
```



```
SELECT * FROM dots;
```



```
SELECT * FROM dots ORDER BY id DESC;
```



```
SELECT * FROM dots WHERE color='black';
```



- ▶ Official documentation on SELECT

## Again, SQL vs CQL semantics

`ex2-orderby-sql.sh`

`ex2-orderby-cql.py`

`ex2-orderby-goodschema.py`

# SELECT

## Order by

- ▶ The order in which partitioned rows are returned depends on the order of the tokens and not on the key values themselves. ([post](#))
- ▶ Clustering column(s) determine the data's sort order only within a partition.
- ▶ ORDER BY not necessary if order is in table definition.

ORDER BY can be used only

- ▶ on clustering columns (i.e. not the partition key)
- ▶ if the partition key is restricted via IN or EQ  
but check paging ...

## ex3-filtering-cql.py

# SELECT

## Where

A WHERE clause can use only the columns that

- ▶ are part of primary key, or
- ▶ have a secondary index on them not discussed so far...

Rationale: if a WHERE clause

- ▶ contains a partition key, Cassandra knows which node to contact: response time is constant wrt the number of nodes.
- ▶ does not contain a partition key, Cassandra contacts each node, and it scans the complete dataset
  - execution time proportional to the total amount of data stored in the referenced table,
  - performance is likely to get worse as you add nodes to the cluster.

in general

ALLOW FILTERING signals poor data modelling

take home message

Rethink your data modelling !

`ex4-index.py`

# DELETE

## Example

```
delete name from students_good where sid=1164220;
```



```
delete name from students_good
```

```
    where uni='P7' and sid=1164220;
```



```
delete name from students_good where uni='P7';
```



```
delete name from students_good;
```



- ▶ all deletions within the same partition are applied atomically and in isolation
- ▶ DELETE can be conditional through an IF clause, but internally Paxos will be used, worsening the performances.
- ▶ Official documentation on DELETE

# BATCH

Way to group together multiple upserts and deletes.

## Example

```
BEGIN BATCH  
modif1  
modif2  
:  
modifn,  
APPLY BATCH;
```

What for ?

- ▶ Saves network round-trips if contains multiple UPDATE
- ▶ Updates in a batch performed in isolation within partitions
- ▶ Eventually all the modifs will take place, or none will
- ▶ Official documentation on BATCH

Batches are *not* analogues of SQL transactions.

# Transactions ?

Cassandra supports a form of lightweight transactions.

One such transaction

- ▶ might block other lightweight transactions from occurring
- ▶ does not stop normal read and write operations from occurring
- ▶ if mixed with normal operations might generate errors

due to the usage of different timestamp systems

How to create a lightweight transaction ?

Using IF conditions in INSERT, UPDATE, DELETE

- ▶ Implementation via Paxos out of scope
- ▶ Guaranteed consistency: linearisability out of scope
- ▶ Further details [here](#) and [here](#)

## Generalized Consensus and Paxos

Paxos made simple

SwiftPaxos: Fast Geo-Replicated State Machines

## Other differences with SQL

- ▶ JOIN:

JOINS must be computed by the application.

- ▶ FOREIGN KEYS:

Referential integrity must be guaranteed by the application.

Keep in mind you can BATCH updates.

- ▶ tokens

# Tokens

ex5-token.py

## How to find which machine is in charge of a token ?

```
$ nodetool [-p port] [-h host] getendpoints keyspace  
table id
```

```
user@host: $ nodetool getendpoints demo dummy 73  
192.168.0.101  
192.168.0.104  
192.168.0.103
```

`ex6-fk-join.py`

`ex6-music.py`

End of Part III

Geo-replicated databases

and

weak consistency

## Weak consistency in Neo4j

- ▶ **read-committed** isolation level
- ▶ Neo4j allows **lost updates**

*To ensure deterministic behaviour [...], it is necessary to explicitly acquire a write lock on the node in question. In Cypher there is no explicit support for this, but it is possible to work around this limitation by writing to a temporary property.*

That's the story.  
Thank you for the attention



Questions?